

Abhiram Dodda

linkedin.com/in/abhiramdodda | abhiramdodda.github.io | abhiramdodda@gmail.com

EDUCATION

Pennsylvania State University Master of Science in Computer Science and Engineering	Aug 2025 – May 2027 (GPA: 3.89/4)
Gokaraju Rangaraju Institute of Engineering and Technology Bachelor of Technology in Computer Science and Engineering	Dec 2021 – Jun 2025 (GPA: 9.38/10.0)

EXPERIENCE

Software Development Intern <i>Platform3Solutions</i>	May 2025 – Jul 2025 Remote
– Architected a self-hosted Retrieval-Augmented Generation (RAG) pipeline by deploying Llama 3 via an Ollama server, utilizing Flask to facilitate semantic search for 500+ concurrent users with a mean query latency of <200ms . – Optimized retrieval precision by 76% by engineering a hierarchical metadata filtering layer within ChromaDB , integrating fine-grained RBAC to ensure 100% data isolation across multi-tenant project environments. – Engineered scalable, containerized Flask REST APIs orchestrated via Docker , managing an automated ingestion pipeline that processed 5,000+ documents per hour with asynchronous local embedding generation.	
Data Science Intern <i>CaptainFresh</i>	Mar 2025 – Jun 2025 Remote
– Executed comprehensive Exploratory Data Analysis (EDA) on 100k+ multi-band spectral observations, utilizing automated statistical profiling to identify significant correlations ($r > 0.82$) between specific spectral bands and pond health indicators. – Engineered a robust feature suite by deriving 12+ spectral indices and applying Principal Component Analysis (PCA) to reduce input dimensionality by 30% while preserving 95% of the variance in the dataset. – Trained and benchmarked various machine learning architectures (Random Forest, XGBoost) to predict pond ecosystem conditions, achieving an R-squared of 0.88 through iterative hyperparameter tuning and 5-fold cross-validation.	
Research Assistant <i>GRIET</i>	Sep 2023 – Mar 2025 <i>Hyderabad, India</i>
– Developed high-precision diagnostic systems ; designed a CNN for dermatological classification (95.27% accuracy) and a plant pathology framework with a 96.82% top-1 accuracy across 38 distinct classes. – Designed a multi-model ensemble architecture for Optical Coherence Tomography (OCT) analysis, fusing VGG16, InceptionV3, and U-Net backends to reach a state-of-the-art accuracy of 98.86% and an AUC-ROC of 0.99 . – Developed a patented Assistive AI system for the visually impaired, integrating a custom YOLO-based object detection model (operating at 30+ FPS) with Tesseract OCR to provide sub-second auditory feedback for obstacle avoidance. – Mitigated severe class imbalance in medical datasets through a hybrid pipeline of synthetic data augmentation and strategic under-sampling , improving minority class F1-scores by 14% .	
Research Intern <i>IIT Hyderabad</i>	May 2024 – Dec 2024 <i>Hyderabad, India</i>
– Quantified semantic stability in Transformers by proposing a semi-intrinsic evaluation metric ; analyzed probability distribution shifts (median $\Delta \approx 0.03$) across 10^4 synonym substitutions to measure embedding robustness . – Fine-tuned BERT-large models for domain-specific sentiment analysis , achieving an 80% F1-score and optimizing hyperparameters to reduce training convergence time by 20% . – Optimized Information Retrieval (IR) systems for judicial cross-referencing, utilizing BM25 and Cosine Similarity to achieve a Mean Reciprocal Rank (MRR) of 0.82 across heterogeneous legal datasets.	

PUBLICATIONS & PATENTS

Publications

- **A Novel Deep-Learning Based Classification of Skin Diseases**
Tabitha Indupalli, Dodda Abhiram*, Rayapuraju Srivatsav, Singamsetty Aashrith, Vudiga Vineeth*
5th IEEE Global Conference for Advancement in Technology (GCAT), 2025
Developed a Deep CNN achieving 95.27% accuracy for multi-class skin disease classification using data augmentation + under-sampling to address imbalance.

- **A Novel Two-Stage Deep Learning Framework for Plant Disease Detection**
Dodda Abhiram, K. Anuradha, V.Srilakshmi, K. Adilakshmi*
 5th IEEE Global Conference for Advancement in Technology (GCAT), 2025
 Proposed a two-stage hierarchical model (species → disease) achieving 96.82% accuracy on PlantVillage dataset.
- **Deep Learning and OCT Imaging: A Novel Ensemble Approach for Eye Disease Diagnosis**
Dodda Abhiram, R. Aruna Flarance, V. Srilakshmi, K. Anuradha*
 4th International Conference on Cognitive & Intelligent Computing (ICCIC) – Springer
 Ensemble of VGG16, InceptionV3, and InceptionResNetV2, achieving 98.86% accuracy in OCT-based eye disease detection.
- **Ensemble Approach for Blood Vessel Segmentation in Retinal Images**
Dodda Abhiram, Y. Ekantha Sai, B. Siddarth Mahesh, K. Anuradha, B. Sankara Babu*
 International Conference on Intelligent Systems & Sustainable Computing (ICISSC) – Springer
 Hybrid U-Net + SegNet with logistic regression refinement, achieving 97.021% segmentation accuracy.

Patent

- **Smart Glasses for Visually Impaired Individuals**
Sreejyothsna Ankam, Manav, Abhiram Dodda, Dinesh Chandra, Nagireddy Padmakshaya
 Patent No. 202541000874 A, Issued Jan 11, 2025
 Wearable assistive device with object detection, obstacle recognition, OCR-based reading, currency detection, and voice-controlled smartphone integration.

PROJECTS

- Distributed ML Inference Engine** | C++, ONNX, CUDA, Consistent Hashing, Multi-threading [Link](#)
- Engineered a high-performance C++17 distributed inference system using **ONNX Runtime** with optional **CUDA acceleration** and automatic CPU fallback for heterogeneous hardware support.
 - Achieved a throughput of **522.64 req/sec** and a **100% success rate** by implementing a **Dynamic Batching** layer (32-request max size) and a thread-safe **LRU Cache** maintaining a **99.7% hit rate**.
 - Architected a robust Gateway utilizing **Consistent Hashing** with 150 virtual nodes for load balancing and **Circuit Breakers** to manage worker health, resulting in a **P99 latency of 164.29ms**.
- Speech to Sign Translation** | Whisper, T5-Transformer, PyTorch, OpenCV, NLP [Link](#)
- Architected a multimodal pipeline integrating **Whisper-Large-v3** for robust speech-to-text and a **fine-tuned T5-Transformer** for linguistic text-to-gloss translation.
 - Achieved a **77% translation accuracy** and a **BLEU score of 0.42**, benchmarking the system against standard ASL datasets.
 - Maintained a sub-second (**1.1s**) average inference time per sequence through **INT8 model quantization** and optimized tokenization streams.

- FarmEase: Two-Stage Diagnostic System** | Tensorflow, EfficientNet, YOLO, Flask [Link](#)
- Developed a dual-stage framework using an **EfficientNet-based localizer** for leaf isolation and a **Stage-2 classifier** for disease identification.
 - Achieved **96.8% accuracy** and a **94.2% macro-F1 score** across 38 classes, addressing high-variance agricultural environments.
 - Implemented **Stochastic Gradient Descent with Restarts (SGDR)** to improve convergence stability and prevent the model from settling in local minima.

ACHIEVEMENTS

- **Merit Endowment Award** (2021-2022), Gokaraju Rangaraju Institute of Engineering and Technology
- **Merit Endowment Award** (2022-2023), Gokaraju Rangaraju Institute of Engineering and Technology
- **Merit Endowment Award** (2025), Gokaraju Rangaraju Institute of Engineering and Technology
- Ranked in **Top 5%** nationally: IIT Madras NPTEL Data Science for Engineers.

SKILLS

Languages: C/C++, Python, Java, JavaScript, SQL, HTML, CSS, Shell Scripting
 Frameworks & Libraries: Flask, Django, VueJS, PyTorch, TensorFlow, JAX, scikit-learn, pandas, NumPy, CUDA
 Big Data & Databases: PySpark, Apache Spark, PostgreSQL, MySQL, NoSQL (MongoDB, Redis), ETL
 Tools: AWS (EC2, S3, Lambda), Docker, Git, GitHub, Linux, Terraform, gRPC, RESTful API, Latex, Tableau